

Weakly Supervised Learning in Computer Vision Tutorial

ECCV 2020




Rodrigo Benenson
Google Zurich




Seong Joon Oh
Naver




Hakan Bilen
University of Edinburgh



AlexNet



ResNet



Large-scale
labeled dataset

Deep
networks

IMAGENET

1M boxes for >3000 categories

[Russakovsky IJCV 15]



118K pixelwise labeled images for
80 object and 91 stuff categories

[Lin ECCV 14, Ceasar CVPR 18]

Open Images



15M boxes on 600 categories
2,8M instance segmentations for
350 categories




[Kuznetsova IJCV 18]

Cityscapes



5K images for 30 categories

[Cordts CVPR 16]



{motorbike, person} {motorbike (point),
person (point)}

{motorbike (b-box),
person (b-box)}


{motorbike (pixel labels),
person (pixel labels)}

1 sec
per class

2.4 sec
per instance

10 sec
per instance

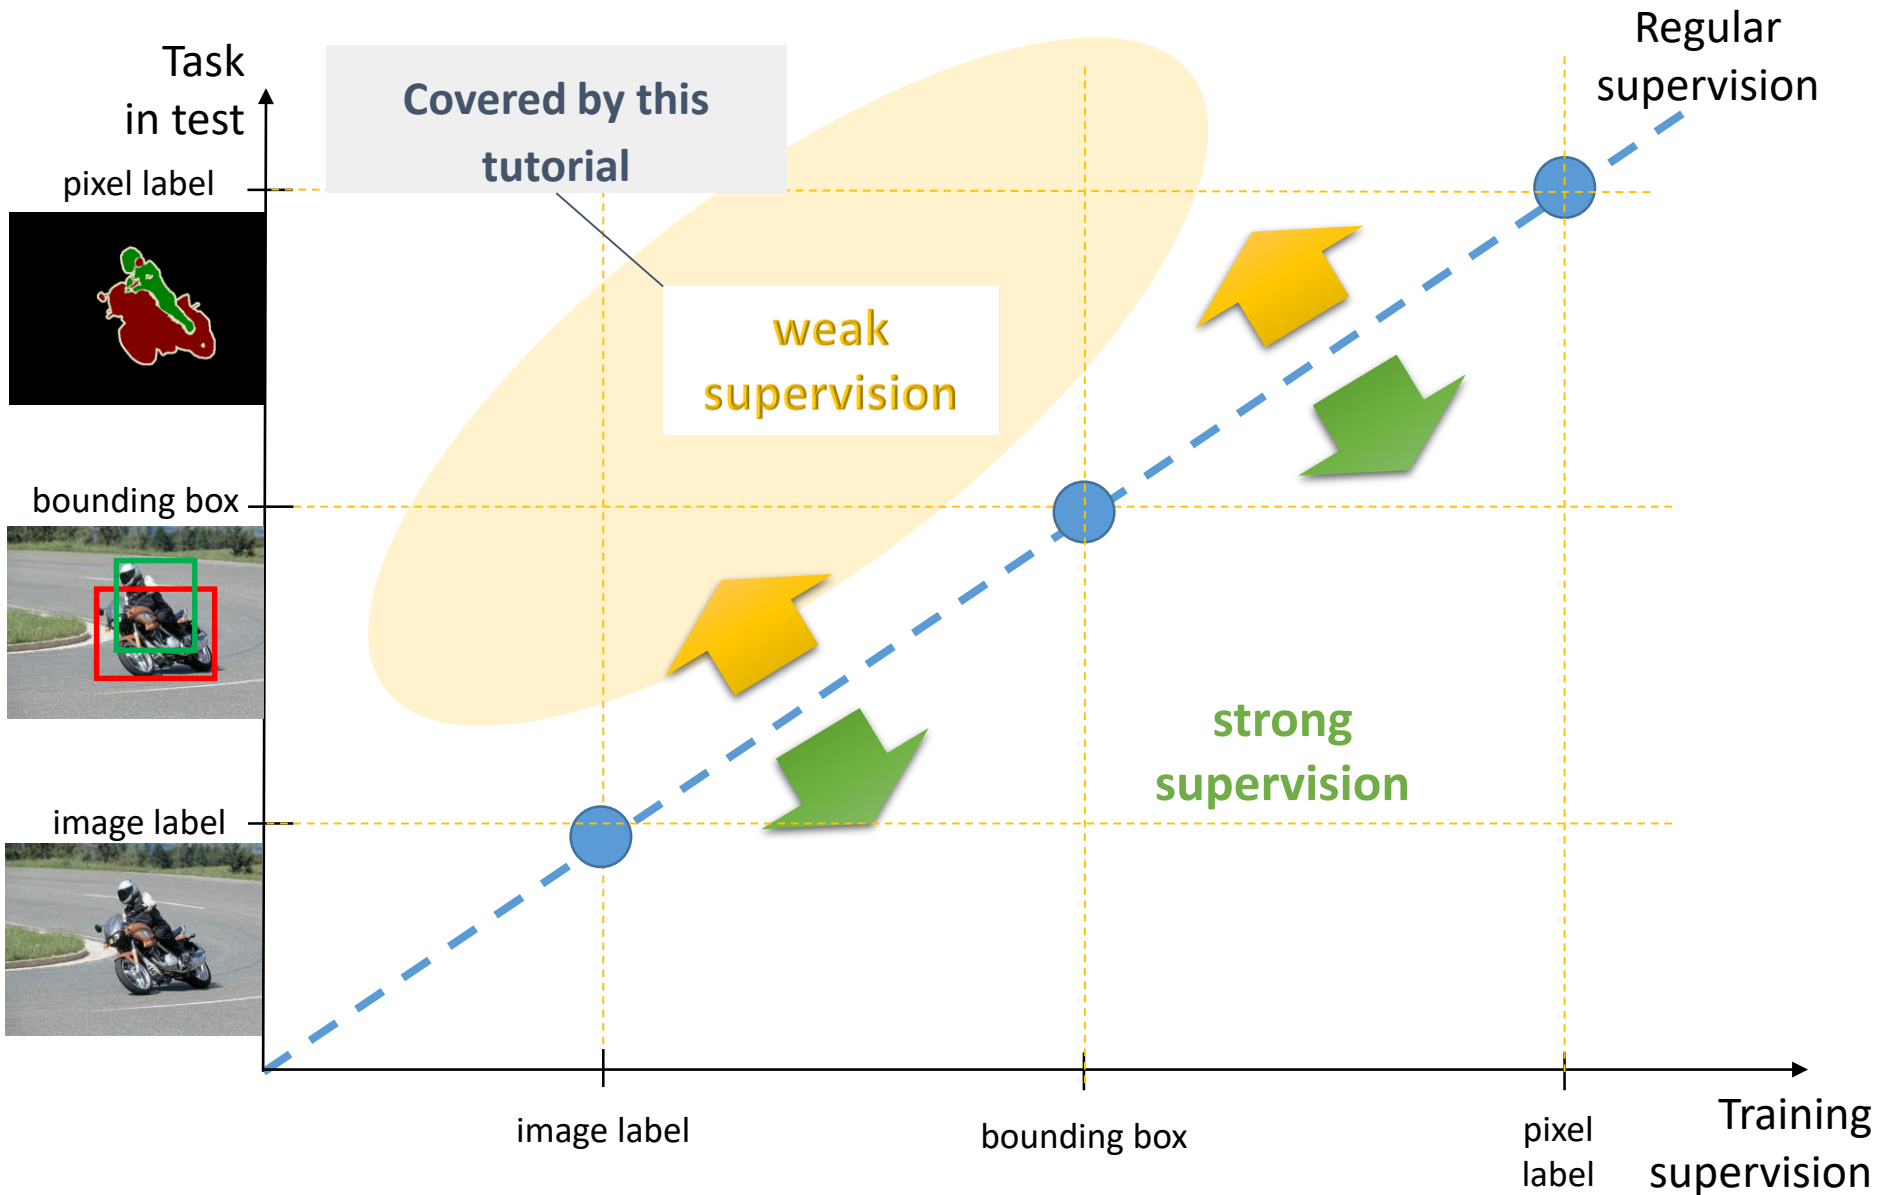
78 sec
per instance



Berman et al., What's the Point: Semantic Segmentation with Point Supervision, ECCV 16

Weak supervision

Lower degree (or cheaper) annotation at train time than the required output at test time





(a) Supervised learning



(b) Weakly supervised learning

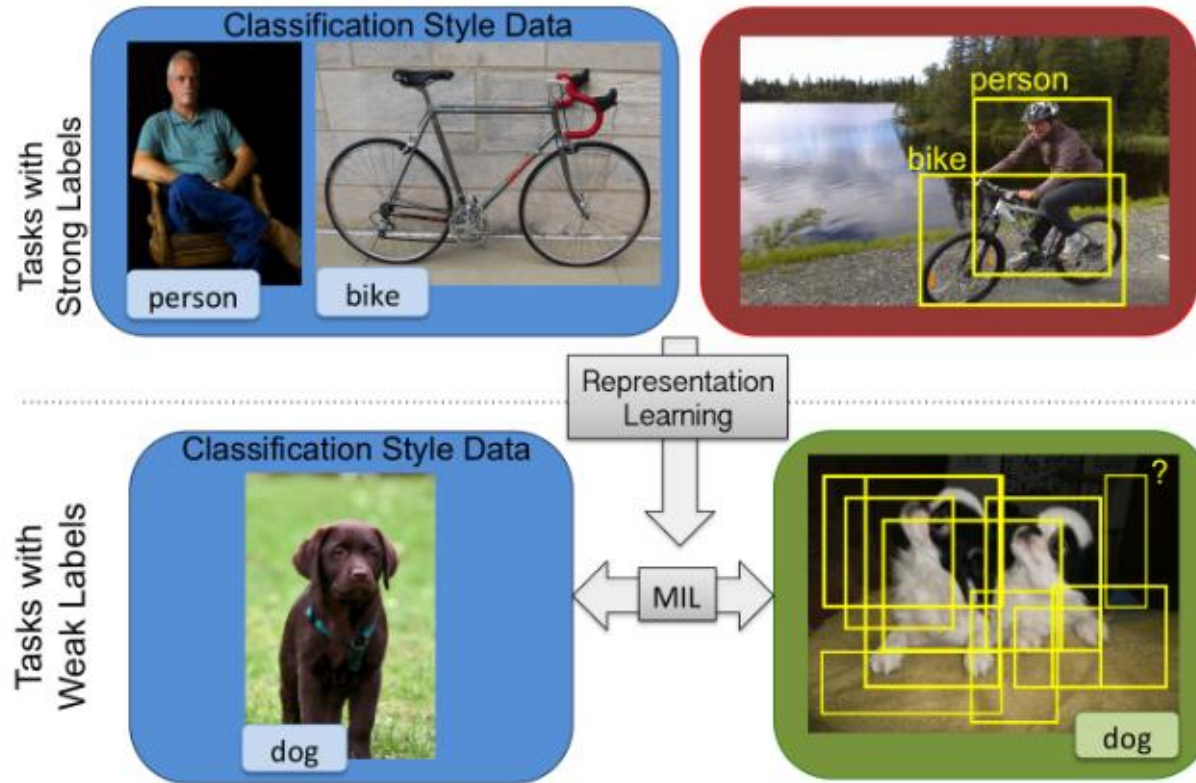


(c) Weakly semi-supervised learning



(d) Semi-supervised learning


- **Semi-supervised**: learns from a set of labeled (bounding box) and unlabeled images
[Wang CVPR 19]
- **Weakly semi-supervised**: learns from a set of labeled (bounding box) and weakly labeled (image label) images
[Yan arXiv 1702.08740]



- Learns detectors for auxiliary categories with regular supervision (bounding box)
- Transfers the knowledge from auxiliary categories to weakly supervised categories

[Hoffman NIPS 14], [Tang CVPR 16]

Example:



Question 1:



Question 2:



Image credits: [Doersch, ICCV 15]

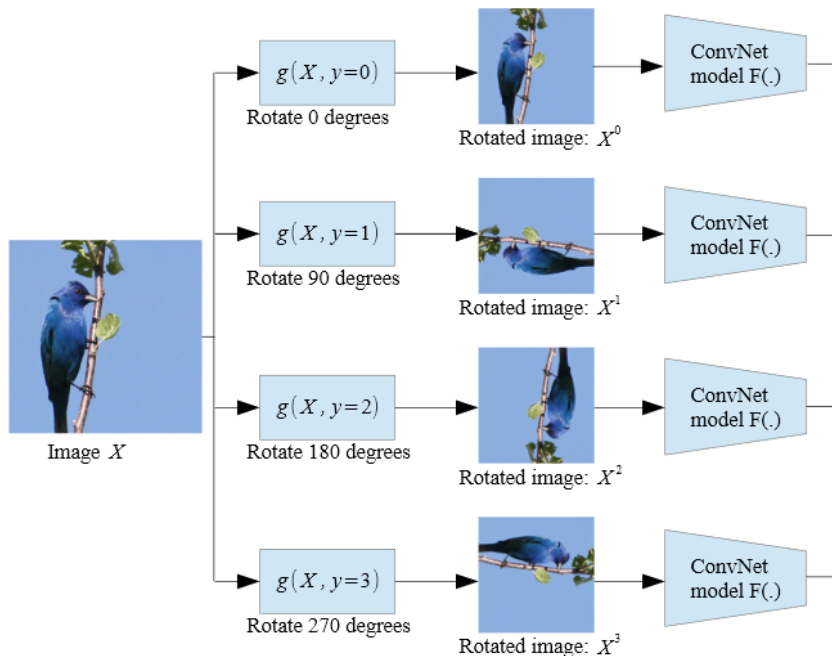


Image credits: [Gidaris ICLR 18]

Self-supervised learning

- Defines a pretext task and learns representations in an unsupervised way
- Transfers pre-trained weights to a downstream task
- Transfer-learning requires regular annotations

Weakly supervised learning

- Typically uses pretrained networks and fine-tunes them by using weak supervision

Both problems reduce dependency to annotations and are complementary!

Generic Object Recognition

Intra-class variations

- Appearance
- Viewpoint
- Scale
- Aspect ratio


Background clutter

Occlusions



Weakly Supervised Learning

Parts vs object



Question: What is a person?

- a) Face
- b) Face + upper body
- c) Face + whole body 😊

Weakly Supervised Learning

Context vs object



Question: What is a train?

- a) Rail
- b) Train + Rail
- c) Train 😊

Weakly Supervised Learning

Multiple vs single instance


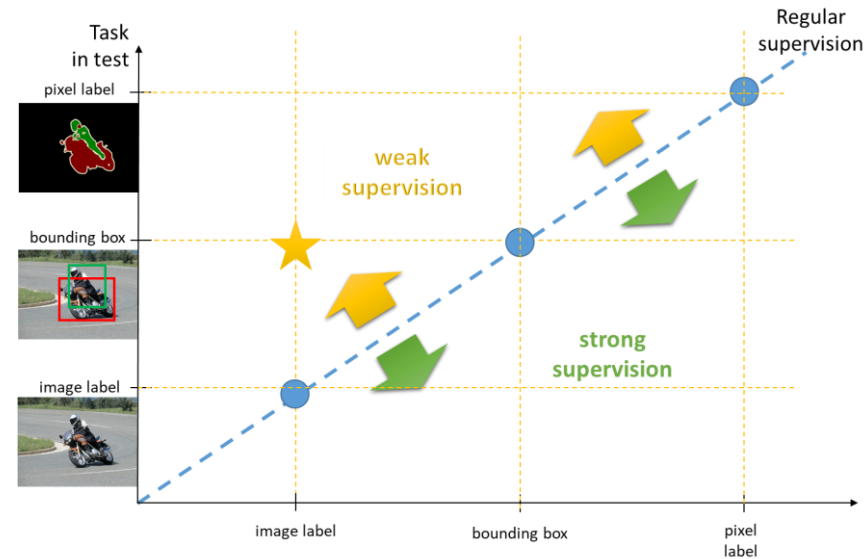


Question: What is a sheep?

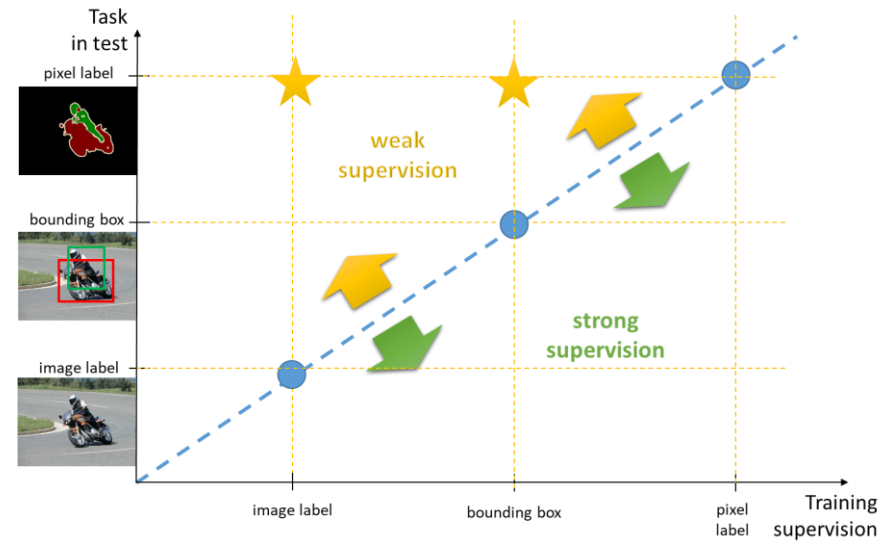
- a) 2 sheeps
- b) 3 sheeps + grass
- c) 1 sheep 😊

Hakan Bilen

- Focuses on object detection with image-level supervision
- Review ideas from 50+ papers
- Provides a framework to the problem
- Discuss challenges and existing solutions



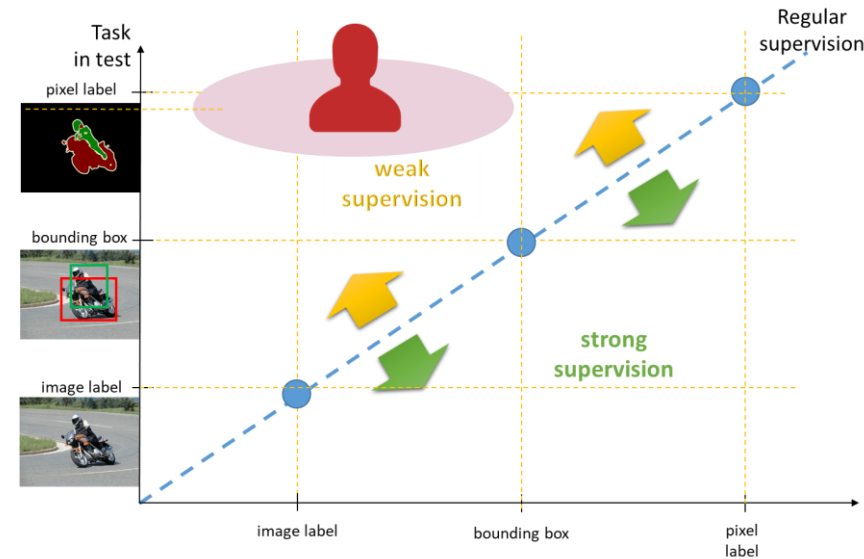
Rodrigo Benenson



- Shares the insights of 40+ papers.
- Provides a framework to read (and write) weakly supervised papers.
- Discusses use of image-labels and box-labels, for semantic and instance segmentation.

Rodrigo Benenson

- Discusses the gap between papers and practice.
- Describes active learning strategies for box and segmentation annotations.
- Presents couple of example papers in the field.



Detection

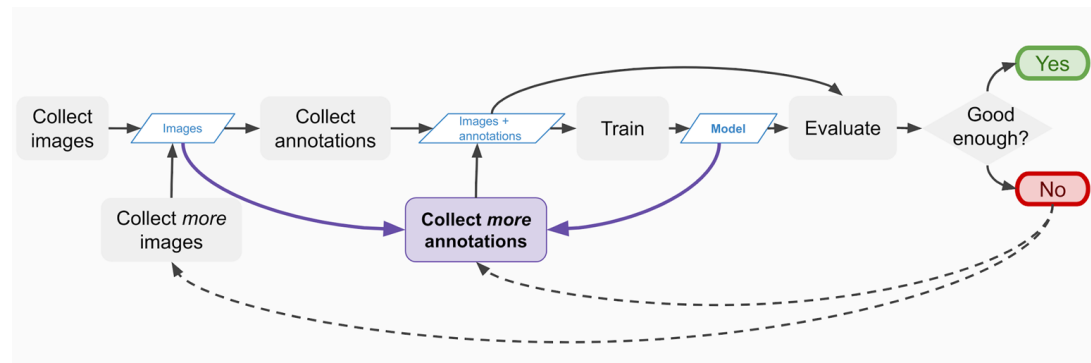


Which boxes to
add?

Semantic labeling


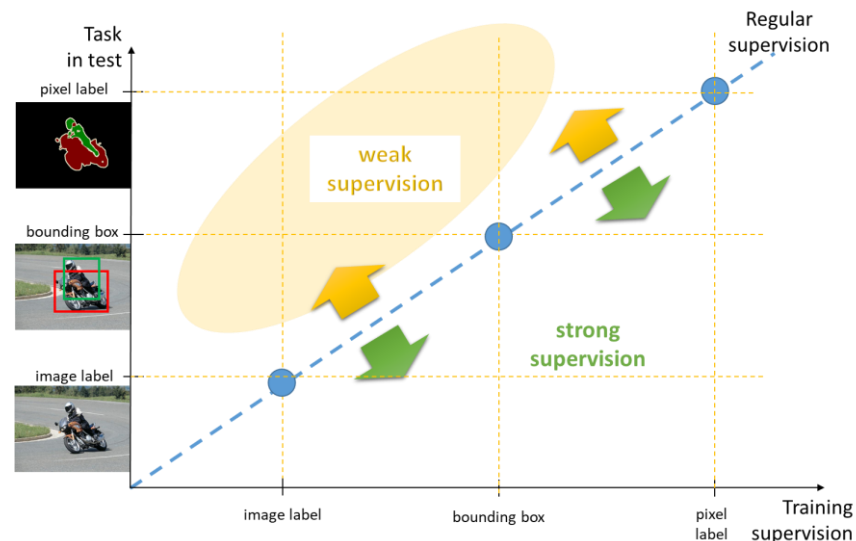


Which pixels to
add?

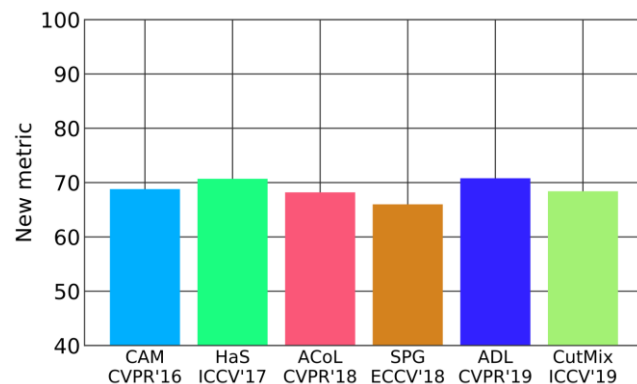


Seong Joon Oh

- Discusses evaluation problem in weakly-supervised learning
- Presents a critical view on the current evaluation practices
- Provides fair evaluation methodologies
- Re-evaluate existing methods in the new benchmark



Improvement according to papers
2016 - 2019



No real improvement under the corrected
evaluation.

Please visit our Q/A sessions

All the videos and slides will be available in our tutorial
webpage:

<https://hbilen.github.io/wsl-eccv20.github.io/>